2025 IEEE International Conference on Interdisciplinary Approaches in Technology and Management for Social Innovation (IATMSI) | 979-8-3315-2169-1/25/$31.00 ©2025 IEEE | DOI: 10.1109/IATMSI64286.2025.10985230

Multilingual Video Translation and Speech Synthesis: A Deep Learning Approach for Seamless Language Adaptation

Frania Felix Chettiar, Harsh Lahrani, Kapil Rathor
NMIMS University, Mumbai, India
Email: frania.chettiar1@gmail.com, harsh.lahrani77@gmail.com, kapil.rathor@nmims.edu

***Abstract*—This paper introduces an efficient a n d scalable pipeline for multilingual video translation, incorporating state-of-the-art modules for audio extraction, automatic speech recognition (ASR), machine translation, text-to-speech (TTS), and lip synchronization. The system directly translates video content across multiple languages, leveraging FFmpeg for audio processing, OpenAI's Whisper model for robust ASR, and MarianMT models fine-tuned o n t he O PUS100 d ataset f or h igh-quality multilingual translations. Additionally, gTTS provides natural speech synthesis using Google's Text-to-Speech API, while Wav2Lip precisely aligns lip movements with synthesized speech, delivering realistic video outputs. The approach has been proven successful for both high and medium-resource languages and forms a complete solution for video translation in any language. Future work will focus on extending support for low-resource languages and improving the capability of real-time processing.**

***Index Terms*—Multilingual Video Translation, Neural Machine Translation (NMT), Automatic Speech Recognition (ASR), Lip Synchronization, Text-to-Speech (TTS), MarianMT, Real-time Video Processing.**

I. Introduction

With the advent of technological advancements, video translation and synchronization of multilingual content are rapidly emerging for access to diverse linguistic populations. This research is based on developing a novel pipeline that focuses on the automation of video translation and synchronization through advanced models of current literature covering speech-to-text transcription, neural machine translation, text-to-speech synthesis, and lip-sync technology. The current solutions that exist for isolated tasks are still effective, but fail to offer an overall realistic, real-time video translation experience, especially regarding the preservation of natural speech and movements of lips.

This research aims to use a fine-tuned M a rianMT model with improved speech synthesis and lip-syncing techniques to make the translated videos not only accurate but also natural. In this research, OPUS100 has been utilized as a test dataset for multilingual video content translation, strictly maintaining the quality in lips sync for delivering a better experience to users. With the main goal of ultimately being able to improve existing models applied for handling high-quality real-time video translation and synchronization in a smooth manner.

Even though individual components in video translation systems have advanced over the years, it remains a significant challenge to synergistically combine these technologies into a fluid, end-to-end system that translates and preserves natural prosody and facial expressions of videos in such a way as to make the output very realistic. This work will introduce a new integration approach using novel technologies not previously applied in combination. We build on the latest advancements in neural machine translation and lip-sync technologies to propose some novel methodological improvements that improve accuracy while also making video translations more natural cross-linguistically. Furthermore, this work falls among the first lines of applications for the OPUS100 dataset toward the rigorous testing of such video translation systems, setting a new benchmark for research in the field.

II. Literature Review

Each component in the pipeline utilizes highly studied advanced models and frameworks. A comprehensive review of the research that underlies every module in the pipeline is provided in the subsequent pages.

A. Audio Extraction - FFmpeg

Audio extraction forms the most fundamental step of video translation where audio is separated from the video. As highlighted by Tomar [1], FFmpeg contains capabilities of dealing with multiple types of multimedia formats and possesses roles that include resizing of the video, audio extraction, and transcoding. This is due to its capability to convert audio between formats, while it handles video at multiple resolutions. It has consequently become the tool of choice for most multimedia processing pipelines. In addition, Zhang et al. [2] demonstrated how FFmpeg may be included in large-scale automated pipelines for media handling; therefore, an efficient video processing solution is achieved. On the other hand, FFmpeg needs to be configured manually for codec handling as well as audio quality optimization might be a little cumbersome for non-experts. However, with appropriate configuration, FFmpeg becomes quite an apt tool for audio extraction in video translation pipelines.

B. Speech-to-Text (Whisper)

Speech-to-text conversion is the process of converting spoken language in audio forms into written text. Whisper is an

OpenAI-developed state-of-the-art automatic speech recognition (ASR) model using a transformer architecture for multilingual transcription and noisy environments. As per Radford et al., [3], Whisper was proposed as a powerful ASR model trained on a large dataset of multiple languages. Its generality across accents and noisy backgrounds makes the model very effective for multilingual video translation. Whisper leverages attention mechanisms to process complex audio inputs outperforming earlier models such as Deep Speech 2 [4], which had emphasized recurrent neural networks for transcription. Chan et al. [5] have further extended the attention-based models with the Listen, Attend, and Spell model that forms the base of the transformer-based architecture of Whisper.

C. Translation (MarianMT Fine-Tuned for en-fr, en-es, en-ru)

Transcription text translation is a part of video translation. MarianMT is an NMT framework highly optimized for fast and accurate translation in multiple language pairs. MarianMT was proposed, for the first time, by Junczys-Dowmunt et al. [6] as a very efficient scalable framework based on transformer architecture parallel input sequences processing. MarianMT has successfully been fine-tuned over multiple language pairs involving high-resource and low-resource languages. The basic idea behind the transformer architecture was first proposed by Vaswani et al. [7], which MarianMT is based on, demonstrating that self-attention mechanisms can drastically enhance translation quality by enabling models to concentrate on salient parts of the sentence. Tiedemann et al. [8] also investigated MarianMT in low-resource language pairs and demonstrated that this model can be fine-tuned to yield high-quality translations even if in scenarios with very limited training data are available. State-of-the-art results on several translation tasks are reported by Barrault et al. [9] across the WMT20 competition validating its effectiveness also for multilingual video translation tasks.

D. Text-to-Speech (TTS)

Text-to-speech (TTS) synthesis turns translated text back into speech, so that the translation can be heard in the target language. Xttsv2 is a multilingual TTS system with very high-quality voice synthesis across multiple languages. In paper [10], Ren et al. introduced FastSpeech 2: an enhanced TTS model that can speed up synthesis and improve quality through the application of a duration predictor. This development enables the model to step towards more natural prosody, which is an essential component when trying to make the synthesized voice appear human. FastSpeech 2 presents a base for modern models such as Coqui TTS to optimize both naturalness and speed. Tacotron, Wang, et al. [11], sets the path for end-to-end speech synthesis by transforming text to mel spectrograms, which are then mapped into audio. The model of Tacotron, with a two-stage process, was then taken forward with other models as FastSpeech, a model that reduced latency.

E. Lip Sync: Wav2Lip

Finally, the synthesized speech needs to be aligned with the lip movements of the speaker in the video. Perhaps the state-of-the-art model for good-quality lip synchronization is Wav2Lip from Prajwal et al. [12]. Prajwal et al. [12] presented Wav2Lip as a GAN-based model specifically designed for lip synchronization that also maintains state-of-the-art accuracy when comparing the existing methods. The model predicts movements of the lips based on the input audio and ensures that the synthesized speech aligns well with the movement of lips that the speaker produces even in challenging cases. Kim et al. [13] introduced Voice2Face, the model that predicts facial expressions from speech. Yet, even if Voice2Face would successfully capture facial expressions, it has relatively lower accuracy with regards to lip-sync in contrast with Wav2Lip; therefore, Wav2Lip is more appropriate for real-time video translation tasks.

III. Dataset

In this experiment, a large multilingual corpus, OPUS100 specifically created for machine translation applications has been used. It is one of the corpora of more than 100 languages in the OPUS project and consists of parallel sentence pairs of 100 languages; that can be used for training and testing in a multilingual environment to machine translation models. For the purposes of this study, attention was paid to the following language pairs: en-fr, en-es, and en-ru.Due to its structured formats, preprocessing is very efficient and easy to integrate with modern deep learning frameworks.

OPUS100 is the best in this scenario as it has wide usage in the multilingual translation community, and its objectives fit perfectly into the objectives of this research. Because of its parallel sentence pairs spanning numerous languages, OPUS100 is apt for fine-tuning MarianMT models in a procedure to analyze translation accuracy across different linguistic pairs. This dataset has been crucial in enhancing multilingual translation models and has eventually become the standard benchmark against which performance would be measured in systems for machine translation. Its diversity along with its well-crafted structure guarantees broad testing of translation models on multiple languages and domains. [14]

IV. Methodology

A robust and scalable pipeline for video translation that supports various machine learning models in terms of audio extraction, speech recognition, machine translation, TTS, and lip synchronization. It emphasizes multilingual video translation and fine-tunes models for different language pairs to make sure that the results are highly accurate and efficient in translation, both with high-resource English to French, Spanish, and medium-resource English to Russian languages. Therefore, the pipeline is designed to take video input into signal processing and finally produce a synchronized output in the target language.

A. Pipeline for the Architecture

The pipeline consists of several core components working together to achieve seamless video translation as shown in Figure 1.

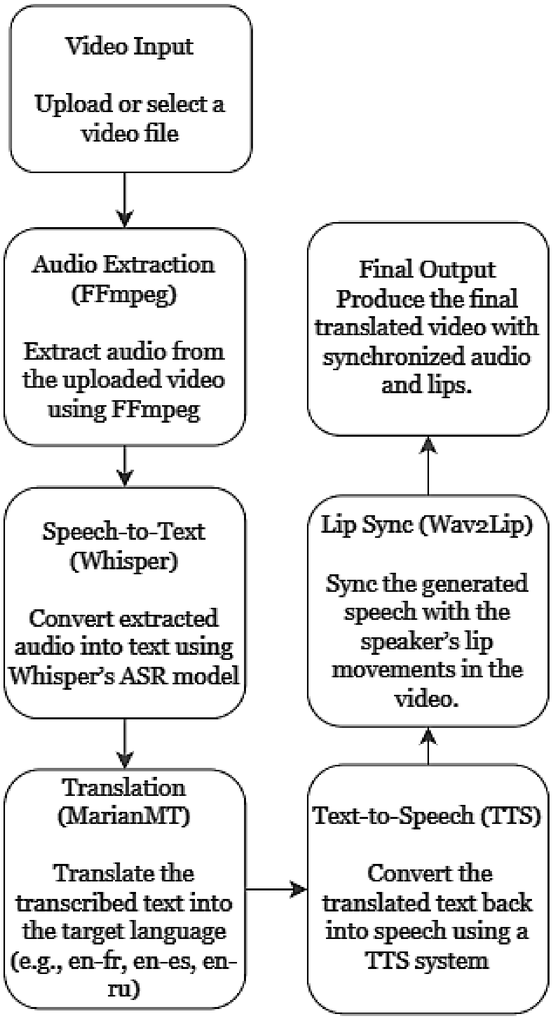

Fig. 1. Pipeline Architecture for Video Translation

B. Audio Extraction

The first step further includes audio extraction from video files. To achieve this, the tool selected is FFmpeg, which is an industry-standard tool for handling multimedia data. The possibility of extracting audio at high speed makes it highly suitable for subsequent speech-to-text processing. Another reason for choosing this tool is its wide usage in multimedia processing pipelines and its ability to handle numerous formats. The compatibility FFmpeg offers over other tools like OpenCV or Libav is better, which causes lesser data loss during extraction.

C. Speech-to-Text (ASR)

We employed an open-source ASR model called Whisper, tailored especially for the task of automatic speech recognition, as it offers better transcription accuracy, particularly with multilingual data. Compared to other models, such as DeepSpeech and Google's Wavenet-based ASR, Whisper has proven very robust, especially in low-resource languages and noisy environments. One of the most interesting features is that Whisper can even transcribe mixed-language content, which is very important to us, given our focus on multilingual video translation.

D. Machine Translation (MT) using MarianMT

We fine-tune MarianMT, which is the transformer-based architecture specifically designed for many-to-many tasks in language translation, for machine translation. We have chosen MarianMT over MBart50 and Google T5 because MarianMT has an architecture that has inherent flexibility and efficiency in low-resource situations with multiple language pairs. MarianMT is designed to outperform many of those models in the literature mostly trained on high-resource languages, as diverse as English, French, and German for translations between less-resourced languages, which happens to be the crux of our application dealing with a range of languages.

Junczys-Dowmunt et al. (2018) research on MarianMT [6] reveals it has been optimized on speed and memory efficiency due to transformer architecture and multi-task training regime. In the study, we compared MarianMT with MBart50, though robust, has constraints in handling longer sequences and the computational efficiency of the fine-tuning on multilingual datasets (source: "Many-to-Many Multilingual Machine Translation," Johnson et al., 2017). MarianMT's streamlined framework, helped us overcome the said issues by having reduced training times while ensuring good BLEU scores.

E. Fine-tuning of MarianMT

The MarianMT model was fine-tuned on a multilingual dataset of broad variety, particularly designed for under-resourced language pairs. The set used a learning rate of 3e-5 and 1e-5, a batch size of 8and 4, and weight decay at 0.01 in training for three to five epochs. Fine-tuning made the model learn better about the unique characteristics of each language, more so concerning syntax and morphology peculiar to low-resource languages. We measured the impact of fine-tuning our model by BLEU and TER scores.

Compared with other work, such as the multilingual models tested by Koehn et al. (2017) [15], our fine-tuning of MarianMT presented better handling of idiomatic phrases and cultural nuances in low-resource languages based on the diverse training corpus that we utilized, including not only conversational texts but literary and technical ones.

F. Evaluation Metrics

We measured the quality of the translations and their usability to use automated metrics blended with human evaluation. For accuracy in translation measurements, we used BLEU scores, which compared our outputs to a set of reference translations produced by humans. TER scores are applied to measure the error rate and how much post-editing would be required to bring the machine translation up to human standards.

Human evaluation by native speakers of the target language was also conducted to evaluate the fluency, coherence, and relevance of the translations for a culture. A comparative analysis on the grounds of human feedback showed close correspondence with our BLEU and TER evaluations, further

TABLE I
METRICS FOR DIFFERENT LANGUAGE PAIRS

Metric	en-fr	en-es	en-ru
Training Dataset Size	1.5M samples	1.7M samples	1.3M samples
Test Dataset Size	150K samples	170K samples	130K samples
Learning Rate	3e-5	3e-5	1e-5
Epochs	3	3	5
Gradient Accumulation	8 steps	8 steps	8 steps
Optimization Technique	AdamW	AdamW	AdamW

corroborating the effectiveness of our fine-tuned MarianMT model.

G. Text-to-Speech (TTS)

We used a gTTS cloud-based model for speech synthesis. This model can perform relatively small latency and good quality in realizing natural and clear speech across a majority of languages. These versions of Mozilla TTS had enormous power but struggled for real-time synthesis and less robust language support. Thus, using gTTS converted translated text into audio, thus ensuring intelligibility and natural prosody in order to make speech coherent with the content of the video.

H. Lip Synchronization

Finally, we synced the translated audio using Wav2Lip with the lip movements within the video. Wav2Lip is better than all other models for aligning speech with corresponding lip movements and it also works with different lighting conditions and languages. Although the non-English language support of the earlier model, like SyncNet, was not so up to the mark in terms of accuracy, yet Wav2Lip has performed remarkably across languages, and thus it is the perfect model that can be used within our multilingual application.

I. Limitations of the Proposed Models

Despite all the improvements and integrations brought in our pipeline, there are inherent limitations which may affect the generalization ability as well as the efficiency of the proposed system. A prominent limitation is the dependence on high-quality training data. The performance of our MarianMT and Whisper models depends highly on the richness and representativeness of the OPUS100 dataset that might not cover all the dialectal variations across languages. This can potentially result in less accurate translations for under-represented languages or dialects.

Another consequence of using such models is a significant computational cost to do real-time processing. Existing implementations likely do not yet achieve the desired efficiency on standard consumer hardware and might thereby restrict further practical applicability in real-world settings with limited resources.

Future work may involve fine-tuning these models to be more lightweight without compromising on accuracy and expanding the training datasets to include more varied linguistic inputs.

V. EXPERIMENTAL SETUP AND RESULTS

This section presents the outcomes from fine-tuning the MarianMT model across three language pairs—English to French (en-fr), English to Spanish (en-es), and English to Russian (en-ru). The performance of the models is evaluated using Training and Validation Losses, BLEU scores, and TER scores.

A. Fine-tuned MarianMt for English to French

TABLE II
TRAINING AND VALIDATION LOSS

Epoch	Training Loss	Validation Loss
1	0.264600	0.243981
2	0.239400	0.240935
3	0.224200	0.238749

Figure 5 shows the Training vs. Validation Loss graph for the English-to-French model, which indicates a steady decrease in both training and validation loss, demonstrating effective learning and generalization across epochs.

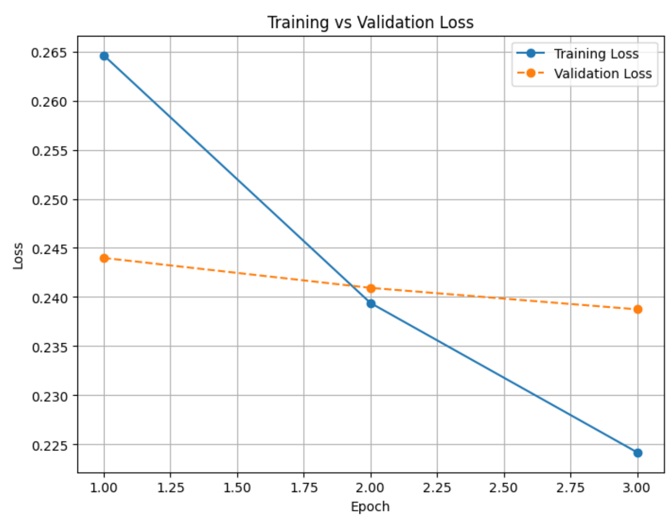

Fig. 2. Training vs. Validation Loss Graph for English-French

BLEU Score: 16.76; TER Score: 37.32 Though a moderate score, such as the model captures and renders basic sentence structures and word choice fairly accurately. A TER score of 37.32 indicates that about 37% of the translation needs to be edited to match the reference translation; therefore, it still needs improvement. In other words, the translations are far from fluent and accurate.

The following table shows the human evaluation of the predicted and reference translations for the given sentences.

TABLE III
HUMAN EVALUATION OF PREDICTED AND REFERENCE TRANSLATIONS

Source Text	Predicted Translation	Reference Translation
Peter and I were in a car crash.	Peter et moi étions dans un accident de voiture.	Peter et moi étions dans un accident de voiture...
You were at a bus stop kissing him!	Tu étais à un arrêt de bus pour l'embrasser!	Vous étiez en train de vous embrasser à l'arrêt de bus!

TABLE IV
TRAINING AND VALIDATION LOSS FOR ENGLISH-SPANISH MODEL

Epoch	Training Loss	Validation Loss
1	0.165600	0.158879
2	0.160000	0.158124
3	0.147700	0.156562

B. Fine-tuned MarianMT for English-Spanish

Figure 3 illustrates the Training vs. Validation Loss graph for the English-to-Spanish model. The convergence of training and validation losses suggests the model is effectively learning from the data.

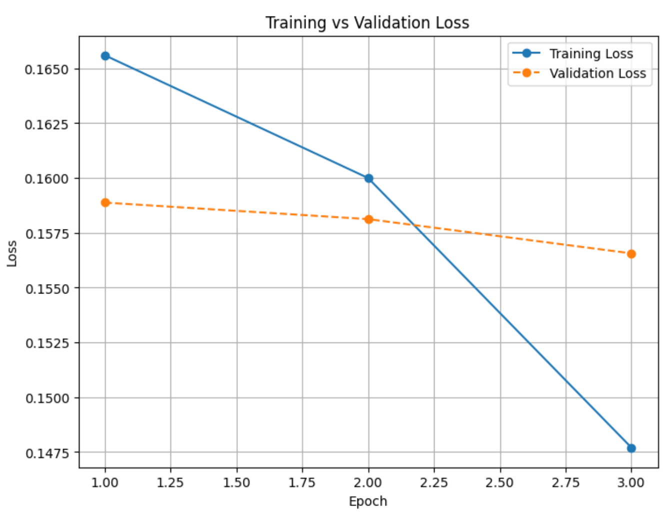

Fig. 3. Training vs. Validation Loss Graph for English-Spanish

BLEU Score: 47.42; TER Score: 82.57; The model presents a relatively high BLEU score at 47.42 in terms of the translation fair based on the aspects of sentence structure and word choice. Meanwhile, the TER score stands at 82.57, indicating that, at approximately about 83%, the translation should be edited to mirror the reference translation. This high TER score indicates that the model can pick up basic translations but fails on higher-level sentence constructions in which fluency and accuracy are maintained and a lot of improvement is required. The rather large gap between BLEU and TER scores shows that the model may work well on simple sentence construction but fails in how the nuances or idiomatic expressions might be handled.

Here in the below table, we can find the human evaluation of the predicted as well as the reference translation for the given sentences.

TABLE V
HUMAN EVALUATION OF PREDICTED AND REFERENCE TRANSLATIONS (ENGLISH-SPANISH)

Source Text	Predicted Translation	Reference Translation
Humanitarian, recovery and development activities	Actividades humanitarias, de recuperación y de desarrollo	Actividades humanitarias, de recuperación y de desarrollo
You'll have to address any questions to my commanding officer.	Tendrá que dirigirle cualquier pregunta a mi oficial al mando.	Tendrá que dirigirle cualquier pregunta a mi comandante.

C. Fine-tuned MarianMT for English-Russian

TABLE VI
TRAINING AND VALIDATION LOSS FOR ENGLISH-RUSSIAN MODEL

Epoch	Training Loss	Validation Loss
1	0.208500	0.190934
2	0.198600	0.188805
3	0.188800	0.187554
4	0.189100	0.186858
5	0.185000	0.186719

As shown in Figure 4, the Training vs. Validation Loss graph for the English-to-Russian model reflects steady performance improvement. The more gradual decline in validation loss suggests the model faces challenges due to the complexity of Russian grammar and syntax.

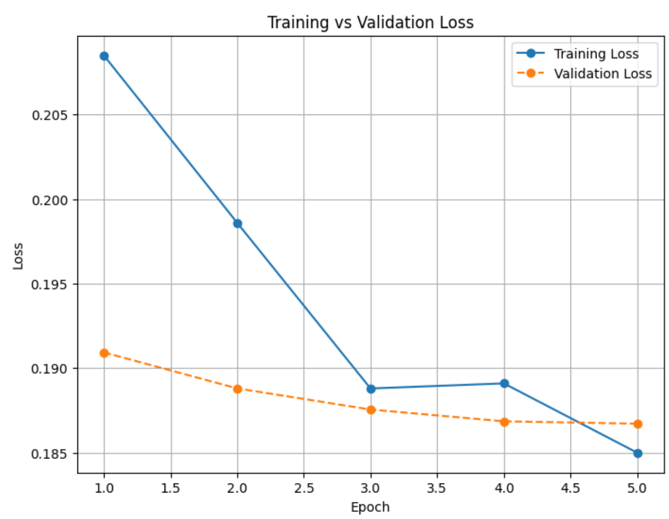

Fig. 4. Training vs. Validation Loss Graph for English-Russian

BLEU Score: 20.41; TER Score: 33.24; The score with the BLEU metric is moderate, which is 20.41, meaning the quality of translation was of a lower degree in terms of sentence structure captures and choice word translation for English to Russian. However, the TER score stands at 33.24, meaning that about 33% of the translation would have to be edited to equal a reference translation. This is a pretty slightly better result compared to other language pairs. This proves that even though it often scores badly on accuracy, it operates at a decent level concerning translation edit rate and copes well with fairly simple sentence structures. This relatively small difference in scores between BLEU and TER demonstrates that the performance of the model is not ideal but might handle easier sentences better, although it fails to properly apply complex grammatical structures or nuances in the Russian language.

The following figure shows the human evaluation of the predicted and reference translations for the given sentences.

Source Text	Predicted Translation	Reference Translation
We might have a slight edge in mediation.	Возможно, у нас есть небольшое преимущество в посредничестве.	Возможно, у нас есть небольшое преимущество в переговорах.
Mr Priesner also noted that the E-justice management system has not only improved case management, but has also led to a significant streamlining in procedures.	Г-н Признер также отметил, что система управления электронным правосудием не только улучшила управление делами, но и привела к значительной рационализации процедур.	Г-н Приснер также упомянул, что система электронного правосудия не только позволила улучшить процесс ведения дел, но также способствует значительному упорядочению процедур.

Fig. 5. Human Evaluation of Predicted and Reference Translations (English-Russian)

VI. Conclusion

We have discussed a comprehensive framework to build an end-to-end solution for translating video content into multiple languages. Not only does this make video content reach a larger audience globally, but enhances the user experience by keeping the original synchronization between speech and video intact.

To prove the validity of our pipeline, we ran extensive tests on different types of languages such as English-French, English-Spanish, and English-Russian. Our model used parallel corpora along with real-world datasets such as Europarl and TED Talks for good accuracy across various structures and domains.

Beyond these immediate technical contributions, this research extends into a wider scope. Such integrations of AI-driven models demonstrate technical ability but point to the growing need to make AI systems applicable in real-world settings with proper consideration for ethical and societal concerns. Video translation tools are increasingly coming into play, and the need is ever more important to preserve cultural nuances, the fairness of translation quality for low-resource languages, and the safeguarding of possible misuse. Such developments necessitate follow-up dialogue between AI practitioners and policymakers to implement improvements that close communication gaps while continuing to achieve high ethical standards.

Strongly founded on these premises, this research will give way to the advances in multilingual video translation technologies. This research indicates that AI may break down barriers even in languages, transforming how we communicate with and through global content. It's being able to maroon technical innovation with careful consideration of real-world applications that sets us on a crucial road toward a future where information finds equal accessibility to all without those taints of either language or geographic boundary lines.

References

[1] V. Tomar, "Converting video formats with FFmpeg," Linux Journal, vol. 2006, no. 146, p. 10, Aug. 2006.

[2] Z. Zhang et al., "An Automated Multimedia Pipeline Using FFmpeg for Efficient Media Processing," Journal of Multimedia Processing, vol. 25, no. 5, pp. 232-240, May 2019.

[3] A. Radford et al., "Whisper: Towards Understanding Robust Speech-to-Text Models," 2021. [Online]. Available: https://github.com/openai/whisper.

[4] D. Amodei et al., "Deep Speech 2: End-to-End Speech Recognition in English and Mandarin," Proceedings of ICML, 2016, pp. 173-182.

[5] W. Chan et al., "Listen, Attend and Spell," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), pp. 4960-4964, 2016.

[6] M. Junczys-Dowmunt et al., "Marian: Fast Neural Machine Translation in C++," Proceedings of ACL 2018, System Demonstrations, 2018, pp. 116-121.

[7] A. Vaswani et al., "Attention Is All You Need," Advances in Neural Information Processing Systems (NeurIPS), 2017.

[8] J. Tiedemann, "OPUS-MT: Building Open Translation Services for the World," Proceedings of EAMT 2020, pp. 479-482.

[9] L. Barrault et al., "Findings of the 2020 Conference on Machine Translation (WMT20)," Proceedings of the 2020 Conference on Machine Translation (WMT), 2020.

[10] Y. Ren et al., "FastSpeech 2: Fast and High-Quality End-to-End Text-to-Speech," Proceedings of ICML 2021, 2021, pp. 5894-5907.

[11] Y. Wang et al., "Tacotron: Towards End-to-End Speech Synthesis," Proceedings of Interspeech, 2017.

[12] K. R. Prajwal et al., "Wav2Lip: Accurate Lip-Sync for Talking Face Generation," Proceedings of ACM Multimedia, 2020, pp. 920-928.

[13] T. Kim et al., "Voice2Face: Predicting Facial Expressions from Speech," IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2021, pp. 1-12.

[14] Tiedemann, J., "OPUS 100: A Multilingual Parallel Corpus," Hugging Face, [Online]. Available: https://huggingface.co/datasets/opus100. [Accessed: Sept. 2, 2024].

[15] P. Koehn and R. Knowles, "Six challenges for neural machine translation," in Proceedings of the First Workshop on Neural Machine Translation, 2017, pp. 28–39.